

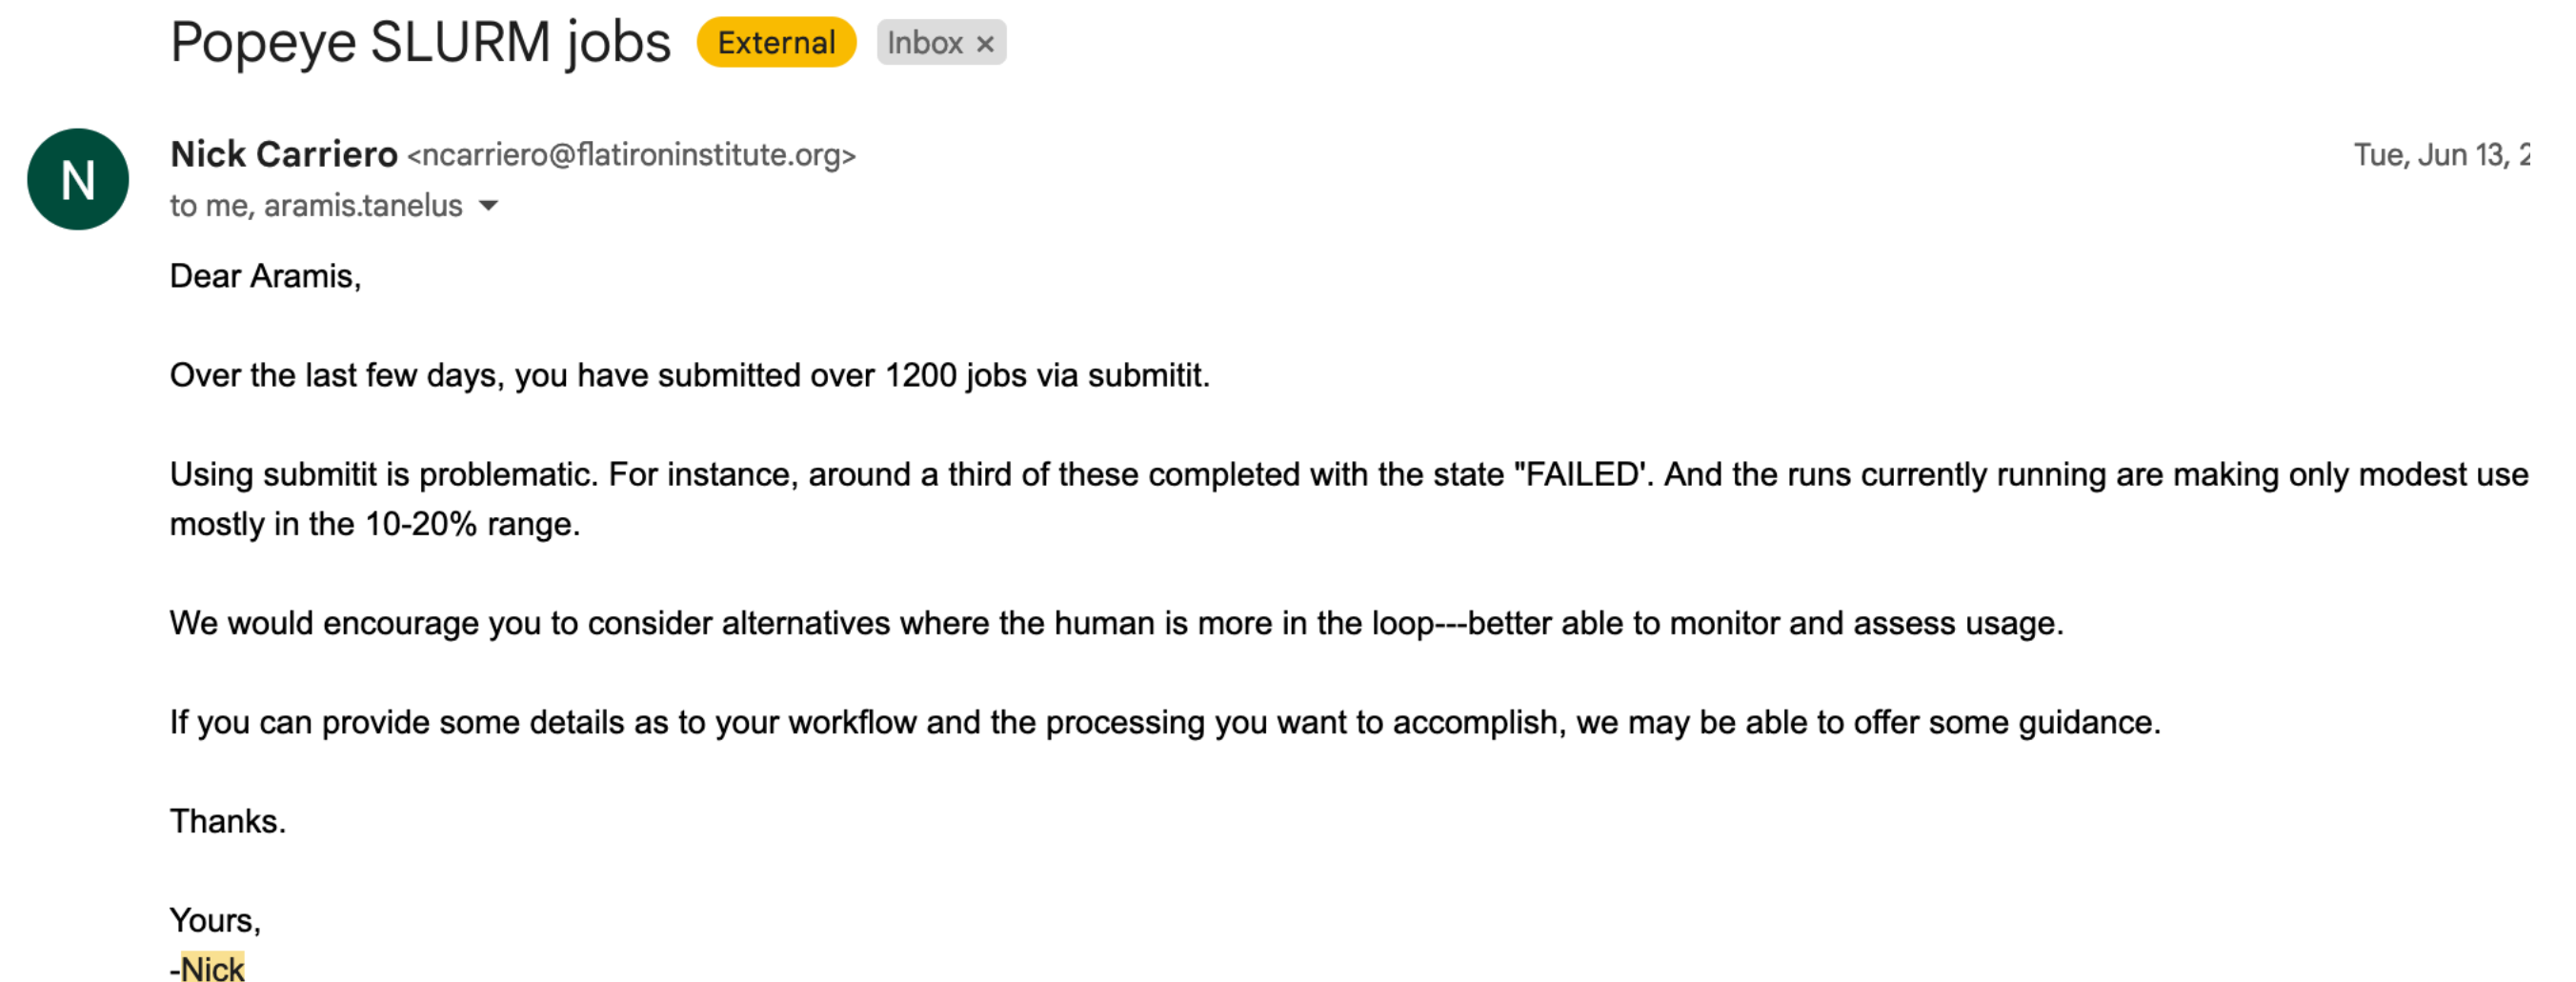
Intro to disBatch

Efficiently schedule and run parallel jobs

2023-08-06

What is disBatch?

- A tool for running many commands across a pool of nodes on the cluster
 - Only requires one resource allocation upfront (saves time in queue)
 - Saves time on initializing shared resources (e.g. copying datasets to local)
 - Saves time spent on reading concerned emails



Running Jobs

Queued Jobs

234112_0

234112_1

234112_2

234112_3

234112_4



How can I use it?

- If your workflow is scripted (you can type ``python my_script.py args...`` and it runs) it is trivial
- If your workflow is not scripted, convert it

Interlude: making Python scripts

- ✨ Live demo ✨ (what could go wrong?)


How can I use it?

- If your workflow is scripted (you can type ``python my_script.py args...`` and it runs) it is trivial
- If your workflow is not scripted, convert it
- Create a text file with commands to run each of your jobs, one per line
 - I usually make a simple python script for this
- Run disBatch on it

Tips

- You should redirect your logs to make errors and output visible
 - Wrap your commands in ``(...) &>path_to_log.log``
 - Optionally, separate stdout and stderr: ``(...) 1>out.log 2>err.log``

My formula:



```
for base_config_path in base_config_dir.iterdir():
    ft_save_path = model_dir / f"{base_config_path.stem}_no_pretrain"
    log_name = f"{base_config_path.stem}_no_pretrain.log"
    command = (
        "python -u -m gerbilizer "
        f"--config {base_config_path} "
        f"--data {finetune_dataset_path} "
        f"--save-path {ft_save_path}"
    )
    command = f"({command}) &> {log_dir / log_name}"
    lines.append(command)
return lines
```

⋮

```
commands = get_commands()
with open("disbatch_script", "w") as ctx:
    ctx.write("\n".join(commands))
```


Tips

- You should redirect your logs to make errors and output visible
 - Wrap your commands in ``(...) &>path_to_log.log``
 - Optionally, separate stdout and stderr: ``(...) 1>out.log 2>err.log``
- You can do everything in one line...
- But if you have to do a lot of setup (loading environments, etc.) it helps to offload everything into a separate script:

```
command = (  
    "source ~/.bashrc && "  
    "source ~/venvs/general/bin/activate && "  
    f"python -u -m gerbilizer --config {base_config_path} --data {pretrain_dataset_path} --save-path {pt_save_path}"  
)  
command = f"({command}) &> {log_dir / log_name}"
```



```
command = f"./train_model.sh {base_config_path} {pretrain_dataset_path} {pt_save_path} {log_dir / log_name}"
```

Scheduling with SLURM

```
[atanelus@rustyamd1 ~]$ sbatch -p gpu --gpus-per-task=1 --mem=32GB  
-c 6 -t 1-0 disBatch -p disbatch_logs/ disbatch_script
```

Slurm Args & Options:

- -p: Partition (gpu, gen, gen, etc.)
- -n: Number of tasks running in parallel
- -c: Num cores (per task)
- -t: Time limit (for everything)
- --mem: RAM (per task)
- --gpus-per-task: bonus points if you can guess this one

disBatch Args & Options:

- -p: Path for saving logs
- -t: max number of tasks running concurrently per node
- Path to script